

# Machine Learning

## Linear Models

Fabio Vandin

October 25<sup>th</sup>, 2021

# Linear Predictors and Affine Functions

Consider  $\mathcal{X} = \mathbb{R}^d$

**“Linear” (affine) functions:**

$$L_d = \{h_{\mathbf{w},b} : \mathbf{w} \in \mathbb{R}^d, b \in \mathbb{R}\}$$

where

$$h_{\mathbf{w},b}(\mathbf{x}) = \langle \mathbf{w}, \mathbf{x} \rangle + b = \left( \sum_{i=1}^d w_i x_i \right) + b$$

**Note:**

- each member of  $L_d$  is a function  $\mathbf{x} \rightarrow \langle \mathbf{w}, \mathbf{x} \rangle + b$
- $b$ : *bias*

# Linear Models

Hypothesis class  $\mathcal{H}$ :  $\phi \circ L_d$ , where  $\phi : \mathbb{R} \rightarrow \mathcal{Y}$

- $h \in \mathcal{H}$  is  $h : \mathbb{R}^d \rightarrow \mathcal{Y}$

$\phi$  depends on the learning problem

## Example

- binary classification,  $\mathcal{Y} = \{-1, 1\} \Rightarrow \phi(z) = \text{sign}(z)$
- regression,  $\mathcal{Y} = \mathbb{R} \Rightarrow \phi(z) = z$

# Equivalent Notation

Given  $\mathbf{x} \in \mathcal{X}$ ,  $\mathbf{w} \in \mathbb{R}^d$ ,  $b \in \mathbb{R}$ , define:

- $\mathbf{w}' = (b, w_1, w_2, \dots, w_d) \in \mathbb{R}^{d+1}$
- $\mathbf{x}' = (1, x_1, x_2, \dots, x_d) \in \mathbb{R}^{d+1}$

Then:

$$h_{\mathbf{w},b}(\mathbf{x}) = \langle \mathbf{w}, \mathbf{x} \rangle + b = \langle \mathbf{w}', \mathbf{x}' \rangle \quad (1)$$

$\Rightarrow$  we will consider bias term as part of  $\mathbf{w}$  and assume  $\mathbf{x} = (1, x_1, x_2, \dots, x_d)$  when needed, with  $h_{\mathbf{w}}(\mathbf{x}) = \langle \mathbf{w}, \mathbf{x} \rangle$

# Linear Regression

$$\mathcal{X} = \mathbb{R}^d, \mathcal{Y} = \mathbb{R}$$

Hypothesis class:

$$\mathcal{H}_{reg} = L_d = \{\mathbf{x} \mapsto \langle \mathbf{w}, \mathbf{x} \rangle + b : \mathbf{w} \in \mathbb{R}^d, b \in \mathbb{R}\}$$

**Note:**  $h \in \mathcal{H}_{reg} : \mathbb{R}^d \rightarrow \mathbb{R}$

Commonly used loss function: *squared-loss*


$$\ell(h, (\mathbf{x}, y)) \stackrel{\text{def}}{=} (h(\mathbf{x}) - y)^2$$

$\Rightarrow$  empirical risk function (training error): *Mean Squared Error*

$$L_S(h) = \frac{1}{m} \sum_{i=1}^m (h(\mathbf{x}_i) - y_i)^2$$

# Linear Regression - Example

$d = 1$



# Least Squares

How to find a ERM hypothesis? *Least Squares* algorithm

Best hypothesis:

$$\arg \min_{\mathbf{w}} L_S(h_{\mathbf{w}}) = \arg \min_{\mathbf{w}} \frac{1}{m} \sum_{i=1}^m (\langle \mathbf{w}, \mathbf{x}_i \rangle - y_i)^2$$

Equivalent formulation:  $\mathbf{w}$  minimizing *Residual Sum of Squares* (RSS), i.e.

$$\arg \min_{\mathbf{w}} \sum_{i=1}^m (\langle \mathbf{w}, \mathbf{x}_i \rangle - y_i)^2$$

## RSS: Matrix Form

Let

$$\mathbf{X} = \begin{bmatrix} \cdots & \mathbf{x}_1 & \cdots \\ \cdots & \mathbf{x}_2 & \cdots \\ \cdots & \vdots & \cdots \\ \cdots & \mathbf{x}_m & \cdots \end{bmatrix}$$

$\mathbf{X}$ : *design matrix*

$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix}$$

$\Rightarrow$  we have that RSS is

$$\sum_{i=1}^m (\langle \mathbf{w}, \mathbf{x}_i \rangle - y_i)^2 = (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w})$$



Want to find  $\mathbf{w}$  that minimizes RSS (*=objective function*):

$$\arg \min_{\mathbf{w}} RSS(\mathbf{w}) = \arg \min_{\mathbf{w}} (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w})$$

How?

Compute gradient  $\frac{\partial RSS(\mathbf{w})}{\partial \mathbf{w}}$  of objective function w.r.t  $\mathbf{w}$  and compare it to 0.

$$\frac{\partial RSS(\mathbf{w})}{\partial \mathbf{w}} = -2\mathbf{X}^T(\mathbf{y} - \mathbf{X}\mathbf{w})$$

Then we need to find  $\mathbf{w}$  such that

$$-2\mathbf{X}^T(\mathbf{y} - \mathbf{X}\mathbf{w}) = 0$$

$$-2\mathbf{X}^T(\mathbf{y} - \mathbf{X}\mathbf{w}) = 0$$

is equivalent to

$$\mathbf{X}^T\mathbf{X}\mathbf{w} = \mathbf{X}^T\mathbf{y}$$

If  $\mathbf{X}^T\mathbf{X}$  is invertible  $\Rightarrow$  solution to ERM problem is:

$$\mathbf{w} = (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{y}$$

# Complexity Considerations

We need to compute

$$(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Algorithm:

- 1 compute  $\mathbf{X}^T \mathbf{X}$ : product of  $(d+1) \times m$  matrix and  $m \times (d+1)$  matrix
- 2 compute  $(\mathbf{X}^T \mathbf{X})^{-1}$  inversion of  $(d+1) \times (d+1)$  matrix
- 3 compute  $(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T$ : product of  $(d+1) \times (d+1)$  matrix and  $(d+1) \times m$  matrix
- 4 compute  $(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$ : product of  $(d+1) \times m$  matrix and  $m \times 1$  matrix

Most expensive operation? Inversion!

$\Rightarrow$  done for  $(d+1) \times (d+1)$  matrix

## $\mathbf{X}^T \mathbf{X}$ not invertible?

How do we get  $\mathbf{w}$  such that

$$\mathbf{X}^T \mathbf{X} \mathbf{w} = \mathbf{X}^T \mathbf{y}$$

if  $\mathbf{X}^T \mathbf{X}$  is not invertible?

Let

$$\mathbf{A} = \mathbf{X}^T \mathbf{X}$$

Let  $\mathbf{A}^+$  be the *generalized inverse* of  $\mathbf{A}$ , i.e.:

$$\mathbf{A} \mathbf{A}^+ \mathbf{A} = \mathbf{A}$$

### Proposition

If  $\mathbf{A} = \mathbf{X}^T \mathbf{X}$  is not invertible, then  $\hat{\mathbf{w}} = \mathbf{A}^+ \mathbf{X}^T \mathbf{y}$  is a solution to  $\mathbf{X}^T \mathbf{X} \mathbf{w} = \mathbf{X}^T \mathbf{y}$ .

# Computing the Generalized Inverse of $\mathbf{A}$

Note  $\mathbf{A} = \mathbf{X}^T \mathbf{X}$  is symmetric  $\Rightarrow$  eigenvalue decomposition of  $\mathbf{A}$ :

$$\mathbf{A} = \mathbf{V} \mathbf{D} \mathbf{V}^T$$

with

- $\mathbf{D}$ : diagonal matrix (entries = eigenvalues of  $\mathbf{A}$ )
- $\mathbf{V}$ : orthonormal matrix ( $\mathbf{V}^T \mathbf{V} = \mathbf{I}_{d \times d}$ )

Define  $\mathbf{D}^+$  diagonal matrix such that:

$$\mathbf{D}_{i,i}^+ = \begin{cases} 0 & \text{if } \mathbf{D}_{i,i} = 0 \\ \frac{1}{\mathbf{D}_{i,i}} & \text{otherwise} \end{cases}$$

Let  $\mathbf{A}^+ = \mathbf{V}\mathbf{D}^+\mathbf{V}^T$

Then

$$\begin{aligned}\mathbf{A}\mathbf{A}^+\mathbf{A} &= \mathbf{V}\mathbf{D}\mathbf{V}^T\mathbf{V}\mathbf{D}^+\mathbf{V}^T\mathbf{V}\mathbf{D}\mathbf{V}^T \\ &= \mathbf{V}\mathbf{D}\mathbf{D}^+\mathbf{D}\mathbf{V}^T \\ &= \mathbf{V}\mathbf{D}\mathbf{V}^T \\ &= \mathbf{A}\end{aligned}$$

$\Rightarrow \mathbf{A}^+$  is a generalized inverse of  $\mathbf{A}$ .

**In practice:** the Moore-Penrose generalized inverse  $\mathbf{A}^\dagger$  of  $\mathbf{A}$  is used, since it can be efficiently computed from the Singular Value Decomposition of  $\mathbf{A}$ .

# Logistic Regression

Learn a function  $h$  from  $\mathbb{R}^d$  to  $[0, 1]$ .

What can this be used for?

Classification!

**Example:** binary classification ( $\mathcal{Y} = \{-1, 1\}$ ) -  $h(\mathbf{x}) = \text{probability}$  that label of  $\mathbf{x}$  is 1.

For simplicity of presentation, we consider binary classification with  $\mathcal{Y} = \{-1, 1\}$ , but similar considerations apply for multiclass classification.


# Logistic Regression: Model

Hypothesis class  $\mathcal{H}$ :  $\phi_{\text{sig}} \circ L_d$ , where  $\phi_{\text{sig}} : \mathbb{R} \rightarrow [0, 1]$  is *sigmoid function*

**Sigmoid function** = “S-shaped” function

For logistic regression, the sigmoid  $\phi_{\text{sig}}$  used is the *logistic regression*:

$$\phi_{\text{sig}}(z) = \frac{1}{1 + e^{-z}}$$





Therefore

$$H_{\text{sig}} = \phi_{\text{sig}} \circ L_d = \{\mathbf{x} \rightarrow \phi_{\text{sig}}(\langle \mathbf{w}, \mathbf{x} \rangle) : \mathbf{w} \in \mathbb{R}^d\}$$

and  $h_{\mathbf{w}}(\mathbf{x}) \in H_{\text{sig}}$  is:

$$h_{\mathbf{w}}(\mathbf{x}) = \frac{1}{1 + e^{-\langle \mathbf{w}, \mathbf{x} \rangle}}$$

Main difference with binary classification with halfspaces: when  $\langle \mathbf{w}, \mathbf{x} \rangle \approx 0$

- halfspace prediction is deterministically 1 or -1
- $\phi_{\text{sig}}(\langle \mathbf{w}, \mathbf{x} \rangle) \approx 1/2 \Rightarrow$  uncertainty in predicted label

# Loss Function

Need to define how bad it is to predict  $h_{\mathbf{w}}(\mathbf{x}) \in [0, 1]$  given that true label is  $y = \pm 1$

## Desiderata

- $h_{\mathbf{w}}(\mathbf{x})$  “large” if  $y = 1$
- $1 - h_{\mathbf{w}}(\mathbf{x})$  “large” if  $y = -1$

Note that

$$\begin{aligned} 1 - h_{\mathbf{w}}(\mathbf{x}) &= 1 - \frac{1}{1 + e^{-\langle \mathbf{w}, \mathbf{x} \rangle}} \\ &= \frac{e^{-\langle \mathbf{w}, \mathbf{x} \rangle}}{1 + e^{-\langle \mathbf{w}, \mathbf{x} \rangle}} \\ &= \frac{1}{1 + e^{\langle \mathbf{w}, \mathbf{x} \rangle}} \end{aligned}$$

Then *reasonable* loss function: increases monotonically with

$$\frac{1}{1 + e^{y\langle \mathbf{w}, \mathbf{x} \rangle}}$$

$\Rightarrow$  *reasonable* loss function: increases monotonically with

$$1 + e^{-y\langle \mathbf{w}, \mathbf{x} \rangle}$$

Loss function for logistic regression:

$$\ell(h_{\mathbf{w}}, (\mathbf{x}, y)) = \log \left( 1 + e^{-y\langle \mathbf{w}, \mathbf{x} \rangle} \right)$$

Therefore, given training set  $S = ((\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m))$  the ERM problem for logistic regression is:

$$\arg \min_{\mathbf{w} \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m \log \left( 1 + e^{-y_i \langle \mathbf{w}, \mathbf{x}_i \rangle} \right)$$

**Notes:** logistic loss function is a *convex function*  $\Rightarrow$  ERM problem can be solved efficiently

Definition may look a bit arbitrary: actually, ERM formulation is the same as the one arising from *Maximum Likelihood Estimation*

# Maximum Likelihood Estimation (MLE) [UML, 24.1]

MLE is a statistical approach for finding the parameters that maximize the joint probability of a given dataset *assuming a specific parametric probability function*.

**Note:** MLE essentially assumes a *generative model* for the data

General approach:

- 1 given training set  $S = ((\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m))$ , assume each  $(\mathbf{x}_i, y_i)$  is i.i.d. from some probability distribution of parameters  $\theta$
- 2 consider  $\mathbb{P}[S|\theta]$  (likelihood of data given parameters)
- 3 *log likelihood*:  $L(S; \theta) = \log(\mathbb{P}[S|\theta])$
- 4 *maximum likelihood estimator*:  $\hat{\theta} = \arg \max_{\theta} L(S; \theta)$

# Logistic Regression and MLE

Assuming  $\mathbf{x}_1, \dots, \mathbf{x}_m$  are fixed, the probability that  $\mathbf{x}_i$  has label  $y_i = 1$  is

$$h_{\mathbf{w}}(\mathbf{x}_i) = \frac{1}{1 + e^{-\langle \mathbf{w}, \mathbf{x}_i \rangle}}$$

while the probability that  $\mathbf{x}_i$  has label  $y_i = -1$  is

$$(1 - h_{\mathbf{w}}(\mathbf{x}_i)) = \frac{1}{1 + e^{\langle \mathbf{w}, \mathbf{x}_i \rangle}}$$

Then the likelihood for training set  $S$  is:

$$\prod_{i=1}^m \left( \frac{1}{1 + e^{-y_i \langle \mathbf{w}, \mathbf{x}_i \rangle}} \right)$$

Therefore the log likelihood is:

$$-\sum_{i=1}^m \log \left( 1 + e^{-y_i \langle \mathbf{w}, \mathbf{x}_i \rangle} \right)$$

And note that the maximum likelihood estimator for  $\mathbf{w}$  is:

$$\arg \max_{\mathbf{w} \in \mathbb{R}^d} - \sum_{i=1}^m \log \left( 1 + e^{-y_i \langle \mathbf{w}, \mathbf{x}_i \rangle} \right) = \arg \min_{\mathbf{w} \in \mathbb{R}^d} \sum_{i=1}^m \log \left( 1 + e^{-y_i \langle \mathbf{w}, \mathbf{x}_i \rangle} \right)$$

$\Rightarrow$  MLE solution is equivalent to ERM solution!


# Linear Classification

$\mathcal{X} = \mathbb{R}^d$ ,  $\mathcal{Y} = \{-1, 1\}$ , 0-1 loss

Hypothesis class = *halfspaces*

$$HS_d = \text{sign} \circ L_d = \{\mathbf{x} \rightarrow \text{sign}(h_{\mathbf{w},b}(\mathbf{x})) : h_{\mathbf{w},b} \in L_d\}$$

**Example:**  $\mathcal{X} = \mathbb{R}^2$





# Finding a Good Hypothesis

Linear classification with hypothesis set  $\mathcal{H}$  = halfspaces.

How do we find a good hypothesis?

Good = minimizes the training error (ERM)

$\Rightarrow$  Perceptron Algorithm (Rosenblatt, 1958)

## Note:

if  $y_i \langle \mathbf{w}, \mathbf{x}_i \rangle > 0$  for all  $i = 1, \dots, m \Rightarrow$  all points are classified correctly by model  $\mathbf{w} \Rightarrow$  *realizability assumption* for training set

**Linearly separable data:** there exists  $\mathbf{w}$  such that:  $y_i \langle \mathbf{w}, \mathbf{x}_i \rangle > 0$

# Perceptron

**Input:** training set  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)$


**initialize**  $\mathbf{w}^{(1)} = (0, \dots, 0)$ ;

**for**  $t = 1, 2, \dots$  **do**

**if**  $\exists i$  s.t.  $y_i \langle \mathbf{w}^{(t)}, \mathbf{x}_i \rangle \leq 0$  **then**  $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} + y_i \mathbf{x}_i$ ;

**else return**  $\mathbf{w}^{(t)}$ ;

Interpretation of update:



Note that:

$$\begin{aligned} y_i \langle \mathbf{w}^{(t+1)}, \mathbf{x}_i \rangle &= y_i \langle \mathbf{w}^{(t)} + y_i \mathbf{x}_i, \mathbf{x}_i \rangle \\ &= y_i \langle \mathbf{w}^{(t)}, \mathbf{x}_i \rangle + \|\mathbf{x}_i\|^2 \end{aligned}$$

$\Rightarrow$  update guides  $\mathbf{w}$  to be “more correct” on  $(\mathbf{x}_i, y_i)$ .

Termination? Depends on the realizability assumption!

# Perceptron with Linearly Separable Data

If data is linearly separable one can prove that the perceptron terminates.

## Proposition

Assume that  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)$  is linearly separable, let:

- $B = \min\{\|\mathbf{w}\| : y_i \langle \mathbf{w}, \mathbf{x}_i \rangle \geq 1 \ \forall i, i = 1, \dots, m, \}$ , and
- $R = \max_i \|\mathbf{x}_i\|$ .

Then the Perceptron algorithm stops after at most  $(RB)^2$  iterations (and when it stops it holds that  $\forall i, i \in \{1, \dots, m\} : y_i \langle \mathbf{w}^{(t)}, \mathbf{x}_i \rangle > 0$ ).

# Perceptron: Notes

- simple to implement
- for separable data
  - termination is guaranteed
  - may require a number of iterations that is exponential in  $d$ ...  
⇒ other approaches (e.g., ILP - Integer Linear Programming) may be better to find ERM solution in such cases
  - potentially multiple solutions, which one is picked depends on starting values
- non separable data?
  - run for some time and keep best solution found up to that point (*pocket algorithm*)

# Bibliography

[UML] Chapter 9:

- no 9.1.1